

```
In [4]: from sklearn.datasets import make_moons
from sklearn.cluster import DBSCAN
from sklearn.neighbors import NearestNeighbors
from scipy.spatial.distance import pdist, squareform
```

```
In [6]: # Generate sample data
X, y = make_moons(n_samples=300, noise=0.1, random_state=0)
# Function to plot data
def plot_data(X, title, ax):
    ax.scatter(X[:, 0], X[:, 1], s=10)
    ax.set_title(title)
    ax.set_xlabel('Feature 1')
    ax.set_ylabel('Feature 2')
# Plot initial data
fig, axs = plt.subplots(2, 2, figsize=(12, 10))
plot_data(X, 'Initial Data', axs[0, 0])

# Compute the distance matrix
dist_matrix = squareform(pdist(X))

# Plot distance matrix
axs[0, 1].imshow(dist_matrix, cmap='hot', interpolation='nearest')
axs[0, 1].set_title('Distance Matrix')
axs[0, 1].set_xlabel('Point Index')
axs[0, 1].set_ylabel('Point Index')

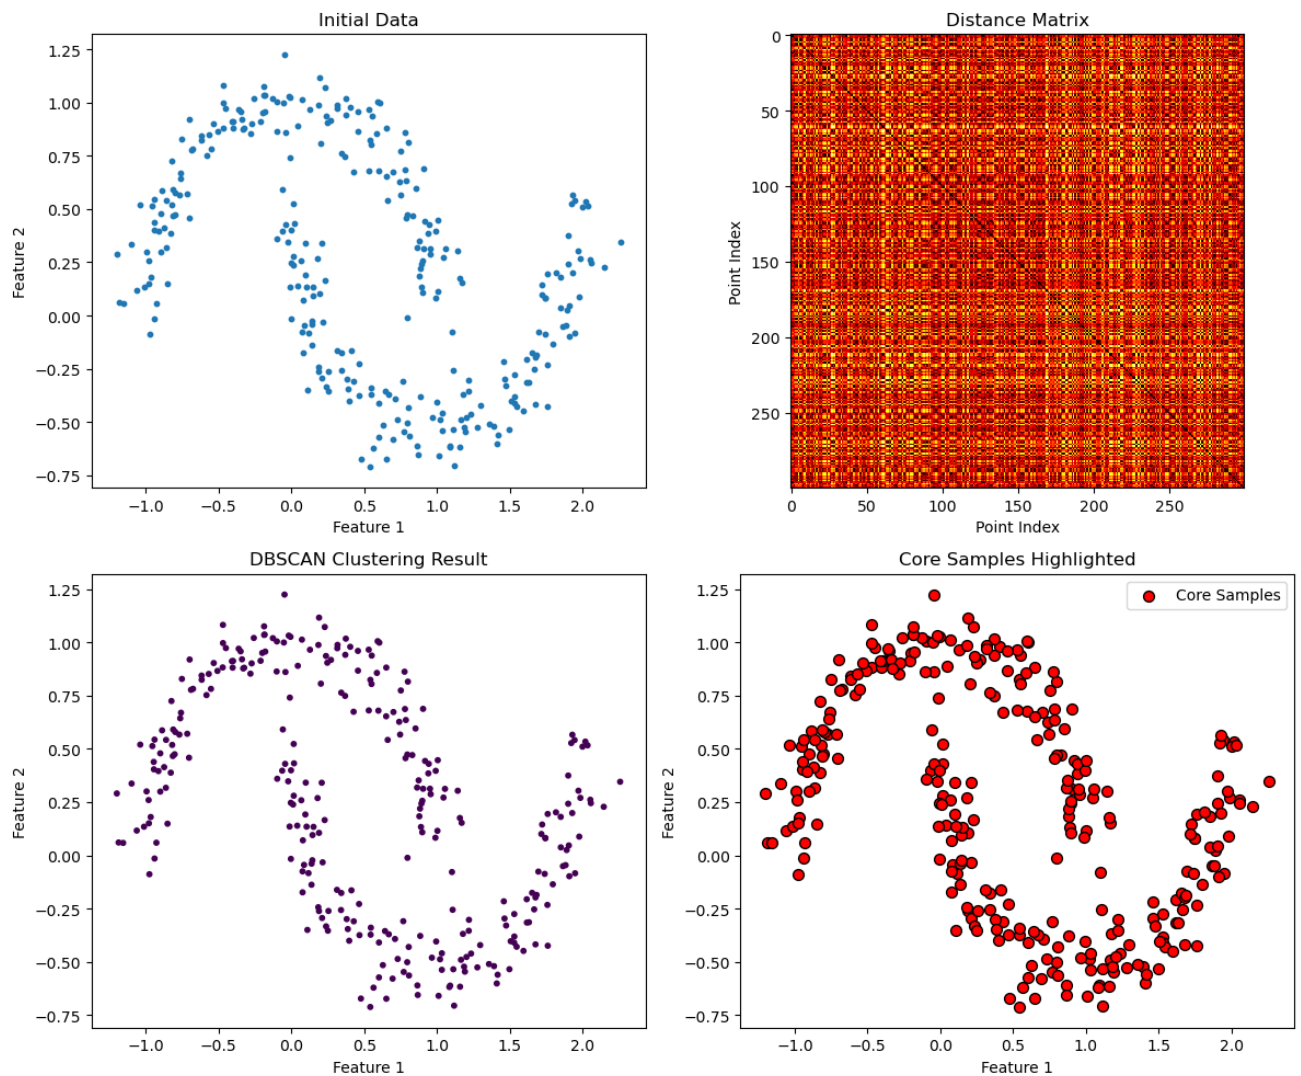
# Apply DBSCAN
eps = 0.2
min_samples = 2
dbscan = DBSCAN(eps=eps, min_samples=min_samples)
labels = dbscan.fit_predict(X)

# Plot clustering result
axs[1, 0].scatter(X[:, 0], X[:, 1], c=labels, cmap='viridis', s=10)
axs[1, 0].set_title('DBSCAN Clustering Result')
axs[1, 0].set_xlabel('Feature 1')
axs[1, 0].set_ylabel('Feature 2')

# Highlight core samples
core_samples_mask = np.zeros_like(labels, dtype=bool)
core_samples_mask[dbscan.core_sample_indices_] = True

axs[1, 1].scatter(X[:, 0], X[:, 1], c=labels, cmap='viridis', s=10)
axs[1, 1].scatter(X[core_samples_mask, 0], X[core_samples_mask, 1], c='red',
axs[1, 1].set_title('Core Samples Highlighted')
axs[1, 1].set_xlabel('Feature 1')
axs[1, 1].set_ylabel('Feature 2')
axs[1, 1].legend()
```

```
plt.tight_layout()  
plt.show()
```



```
In [20]: # load the dataset  
df = pd.read_csv('moons_data.csv')  
# display the first few rows of the dt  
print(df)
```

	x1	x2	y
0	0.771744	-0.548086	1
1	0.189416	-0.261982	1
2	0.918359	0.443277	0
3	1.021213	-0.488523	1
4	1.178442	-0.369193	1
...
295	1.531153	-0.275921	1
296	0.184710	-0.243903	1
297	-0.703622	0.458507	0
298	0.592696	1.006177	0
299	0.104182	0.134738	1

[300 rows x 3 columns]

```
In [9]: # Import libraries
import os
import sys
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings

from scipy.integrate import odeint
from IPython.display import Image
from statistics import mode
from scipy.stats import pearsonr, spearmanr
from sklearn.feature_selection import mutual_info_regression
import statsmodels.api as sm

# Suppress warnings
warnings.filterwarnings('ignore')

# Set number of decimals for np print options
np.set_printoptions(precision=3)

# Set the current working directory
os.chdir(sys.path[0])
```

```
In [51]: # Extract features and labels
X1= df.iloc[:,0]
X2= df.iloc[:,1]
X = df.iloc[:, :2].values
y_true = df.iloc[:,-1]
```

```
In [52]: X1
```

```
Out[52]: 0      0.771744
         1      0.189416
         2      0.918359
         3      1.021213
         4      1.178442
         ...
        295     1.531153
        296     0.184710
        297    -0.703622
        298     0.592696
        299     0.104182
        Name: x1, Length: 300, dtype: float64
```

```
In [53]: X2
```

```
Out[53]: 0      -0.548086
         1      -0.261982
         2       0.443277
         3      -0.488523
         4      -0.369193
         ...
        295    -0.275921
        296    -0.243903
        297     0.458507
        298     1.006177
        299     0.134738
        Name: x2, Length: 300, dtype: float64
```

```
In [35]: y_true
```

```
Out[35]: 0      1
         1      1
         2      0
         3      1
         4      1
         ..
        295     1
        296     1
        297     0
        298     0
        299     1
        Name: y, Length: 300, dtype: int64
```

```
In [36]: # Perform agglomerative clustering
        from sklearn.cluster import AgglomerativeClustering
```

```
In [54]: clustering = AgglomerativeClustering(n_clusters = 7).fit(X)
```

```
In [55]: #print labels
```

```
clustering.labels_
```

```
Out[55]: array([0, 1, 6, 0, 0, 5, 0, 0, 4, 2, 3, 6, 1, 0, 6, 2, 3, 2, 0, 1, 1, 2,
 5, 5, 0, 5, 1, 3, 0, 3, 2, 4, 6, 5, 0, 1, 1, 1, 3, 6, 0, 2, 2, 4,
 4, 5, 0, 4, 6, 4, 2, 0, 1, 2, 4, 3, 1, 6, 3, 1, 5, 5, 4, 3, 3, 0,
 3, 5, 6, 4, 1, 0, 3, 6, 5, 0, 4, 2, 1, 5, 5, 1, 3, 0, 0, 2, 0, 0,
 0, 0, 3, 4, 3, 1, 1, 1, 2, 4, 1, 6, 3, 0, 5, 1, 4, 0, 0, 0, 3, 0,
 1, 1, 4, 2, 0, 4, 1, 3, 5, 6, 4, 5, 6, 4, 2, 2, 1, 1, 1, 6, 2, 1,
 4, 2, 5, 0, 3, 3, 4, 3, 6, 2, 0, 6, 1, 5, 6, 4, 2, 3, 0, 6, 0, 0,
 3, 2, 4, 6, 0, 2, 1, 2, 6, 1, 4, 4, 6, 1, 4, 3, 0, 5, 0, 4, 1, 0,
 5, 0, 0, 4, 2, 5, 3, 0, 5, 2, 0, 2, 0, 2, 0, 1, 4, 1, 5, 0, 2, 3,
 2, 3, 0, 0, 4, 4, 4, 2, 3, 5, 0, 6, 0, 5, 5, 2, 1, 2, 0, 5, 3, 3,
 1, 6, 2, 1, 2, 1, 5, 5, 3, 3, 5, 0, 2, 2, 3, 6, 5, 2, 4, 4, 1, 0,
 4, 6, 5, 3, 1, 4, 2, 5, 2, 6, 3, 0, 4, 5, 4, 0, 5, 4, 1, 1, 1, 0,
 0, 6, 4, 5, 1, 3, 3, 5, 3, 1, 2, 0, 0, 5, 3, 4, 6, 1, 3, 1, 3, 4,
 1, 2, 2, 2, 0, 3, 0, 6, 1, 0, 1, 5, 4, 1])
```

```
In [56]: # Compare true labels with the ones from clustering
y_pred = clustering.labels_
result = pd.DataFrame(np.transpose(np.vstack((y_true, y_pred))), columns= ['
result.iloc[:20,:]
```

Out[56]:

	y_true	y_label
0	1	0
1	1	1
2	0	6
3	1	0
4	1	0
5	0	5
6	1	0
7	1	0
8	0	4
9	0	2
10	1	3
11	0	6
12	1	1
13	1	0
14	0	6
15	0	2
16	1	3
17	0	2
18	1	0
19	1	1

	y_true	y_label
0	1	0
1	1	1
2	0	6
3	1	0
4	1	0
5	0	5
6	1	0
7	1	0
8	0	4
9	0	2
10	1	3
11	0	6
12	1	1
13	1	0
14	0	6
15	0	2
16	1	3
17	0	2
18	1	0
19	1	1

```
In [41]: from sklearn.datasets import make_moons
from sklearn.cluster import DBSCAN
from sklearn.neighbors import NearestNeighbors
from scipy.spatial.distance import pdist, squareform
```

```
In [63]: dbscan = DBSCAN(eps=0.3, min_samples=5)
cluster = dbscan.fit_predict(X)
```


```
In [66]: # Add the predicted cluster labels to the DataFrame for comparison
result_dbscan = pd.DataFrame({
    'X1': X1,
    'X2': X2,
```

```
'y_true': y_true,
'y_label_dbscan': y_pred
})
```

```
In [68]: # Show the first 20 rows of the results
print(result_dbscan.iloc[:20, :])

# Visualize the clustering results
plt.scatter(X1, X2, c=y_pred, cmap='viridis', edgecolor = 'k', s = 50)
plt.xlabel('X1')
plt.ylabel('X2')
plt.title('DBSCAN Clustering')
plt.show()
```

	X1	X2	y_true	y_label_dbscan
0	0.771744	-0.548086	1	0
1	0.189416	-0.261982	1	1
2	0.918359	0.443277	0	6
3	1.021213	-0.488523	1	0
4	1.178442	-0.369193	1	0
5	-1.062691	0.116202	0	5
6	1.234508	-0.461589	1	0
7	0.711839	-0.393012	1	0
8	0.542489	0.824464	0	4
9	-0.752471	0.828908	0	2
10	1.901708	0.242421	1	3
11	0.904850	0.312769	0	6
12	0.210737	-0.294526	1	1
13	1.411097	-0.602078	1	0
14	0.956856	0.287197	0	6
15	-0.099896	1.005161	0	2
16	1.744860	0.080245	1	3
17	-0.450309	0.973776	0	2
18	1.090810	-0.612312	1	0
19	0.095579	0.191930	1	1



```
In [69]: # Import KMeans from sklearn
         from sklearn.cluster import KMeans
```

```
In [70]: kmeans = KMeans(n_clusters=7, random_state=42)
```

```
In [71]: y_pred_kmeans = kmeans.fit_predict(X)
```

```
In [72]: print(y_pred_kmeans)
```

```
[5 4 1 5 5 6 5 5 3 0 2 1 4 5 1 0 2 0 5 4 4 0 6 6 5 6 4 2 5 2 0 3 1 6 5 0 4
 4 2 1 5 0 0 1 3 6 5 3 1 3 0 5 4 0 3 2 4 1 2 4 6 6 3 2 2 5 2 6 1 3 4 5 2 1
 6 2 1 0 4 6 6 4 2 5 5 0 5 5 2 5 2 3 2 4 4 4 0 3 4 1 2 2 6 4 3 5 5 5 2 5 4
 4 3 0 5 3 4 2 6 1 1 6 1 3 0 0 4 4 4 1 0 4 3 0 6 5 2 2 3 2 1 0 5 1 4 6 1 3
 0 2 5 1 2 5 2 0 3 1 5 0 4 0 1 4 3 3 1 4 3 2 5 6 5 3 4 5 6 5 2 3 0 6 2 5 6
 0 5 0 5 0 5 4 1 4 6 5 0 2 0 2 2 5 3 3 3 0 2 6 5 1 5 6 6 0 4 0 5 6 2 2 4 1
 0 4 0 4 6 6 2 2 6 2 0 0 2 1 6 0 1 3 4 5 3 5 6 2 4 1 0 6 0 1 2 5 3 6 3 2 6
 1 4 4 0 2 5 1 3 6 4 2 2 6 2 4 0 5 2 6 2 3 1 4 2 4 2 3 4 0 0 0 5 2 5 1 4 2
 4 6 3 4]
```

```
In [73]: # Add the predicted cluster labels to the DataFrame for comparison
         result_kmeans = pd.DataFrame({
             'X1': X1,
```




```
'X2': X2,
'y_true': y_true,
'y_label_kmeans': y_pred_kmeans
})
```

```
In [74]: # Show the first 20 rows of the results
print(result_kmeans.iloc[:20, :])
```

	X1	X2	y_true	y_label_kmeans
0	0.771744	-0.548086	1	5
1	0.189416	-0.261982	1	4
2	0.918359	0.443277	0	1
3	1.021213	-0.488523	1	5
4	1.178442	-0.369193	1	5
5	-1.062691	0.116202	0	6
6	1.234508	-0.461589	1	5
7	0.711839	-0.393012	1	5
8	0.542489	0.824464	0	3
9	-0.752471	0.828908	0	0
10	1.901708	0.242421	1	2
11	0.904850	0.312769	0	1
12	0.210737	-0.294526	1	4
13	1.411097	-0.602078	1	5
14	0.956856	0.287197	0	1
15	-0.099896	1.005161	0	0
16	1.744860	0.080245	1	2
17	-0.450309	0.973776	0	0
18	1.090810	-0.612312	1	5
19	0.095579	0.191930	1	4

```
In [75]: plt.scatter(X1, X2, c=y_pred_kmeans, cmap='viridis')
plt.xlabel('X1')
plt.ylabel('X2')
plt.title('K-Means Clustering')
plt.show()
```



In []: